

# Kernel Density Estimation

## Assignment 1

Lucas Støjko Andersen

University of Copenhagen

November 8, 2022

# Kernel Density Estimation

## Kernels

Let  $K : \mathbb{R} \longrightarrow [0, \infty)$  with properties

$$K(x) = K(-x), \quad \forall x \in \mathbb{R} \quad (1)$$

$$1 = \int_{\mathbb{R}} K(x) dx \quad (2)$$

and  $X_1, X_2, \dots, X_n$  be random variables with density  $f$ . Then

$$\hat{f}(x) = \frac{1}{n\lambda} \sum_{i=1}^n K\left(\frac{x - x_i}{\lambda}\right) \quad (3)$$

is the kernel density estimate of  $f$  with bandwidth  $\lambda > 0$  and kernel  $K$ .

# Implementation of Kernel Density Estimate

```
R_dens_for <- function(x, p, kernel, bandwidth) {  
  m <- length(p)  
  n <- length(x)  
  result <- numeric(m)  
  for(i in 1:m) {  
    for(j in 1:n) {  
      result[i] <- result[i] + kernel((p[i] - x[j]) / bandwidth)  
    }  
  }  
  result / (n * bandwidth)  
}
```

# Testing the Implementation

## Epanechnikov Kernel

Use the kernel density estimate on the *epanechnikov* kernel given by

$$K(x) = \frac{3}{4}(1 - x^2)1_{[0,1]}(x). \quad (4)$$

Implemented in R as

```
e_kernel <- function(x) {  
  0.75 * (1 - x^2) * (abs(x) <= 1)  
}
```


# The Test Case

## Gaussian Mixture

```
x <- c(rnorm(200, -1, 0.1),  
      rnorm(200, -0.5, 0.1),  
      rnorm(200, 0, 0.3),  
      rnorm(200, 0.5, 0.7),  
      rnorm(200, 1, 1))  
  
funky_density <- function(x) {  
  dnorm(x, -1, 0.1) / 5 +  
  dnorm(x, -0.5, 0.1) / 5 +  
  dnorm(x, 0, 0.3) / 5 +  
  dnorm(x, 0.5, 0.7) / 5 +  
  dnorm(x, 1, 1) / 5  
}
```

# Plotting the Densities

With  $\lambda = 0.1$  we plot the estimated density together with R's default density implementation and the true density.



# Bandwidth Selection

## Leave-One-Out Cross Validation

We pursue bandwidth selection through leave-one-out cross validation. Let

$$\hat{f}_{\lambda}^{-i} = \frac{1}{(n-1)\lambda} \sum_{j \neq i}^n K\left(\frac{x_i - x_j}{\lambda}\right). \quad (5)$$

We then select the bandwidth

$$\hat{\lambda}_{\text{CV}} = \arg \max_{\lambda > 0} \sum_{i=1}^n \log \hat{f}_{\lambda}^{-i}. \quad (6)$$

# Implementation of LOOCV Bandwidth Selection

## Bandwidth Selection and Oracle Bandwidth as Comparison

```
bw_cv_R2 <- function(x, kernel, max_bw = 2) {  
  cv_func <- function(l) {  
    if(l < .Machine$double.eps) Inf  
    n <- length(x)  
    K <- numeric(n)  
    for(i in 1:n) {  
      K[i] <- sum(kernel((x[i] - x[-i]) / l))  
    }  
    cv <- sum(log(K[K > .Machine$double.eps]))  
    n * log((n - 1) * l) - cv  
  }  
  suppressWarnings(optimize(cv_func, c(0, max_bw)))$minimum  
}  
  
bw_oracle <- function(x, kernel) {  
  n <- length(x)  
  K <- integrate(function(x) kernel(x)^2, -Inf, Inf)$value  
  sigma2 <- integrate(function(x) kernel(x) * x^2, -Inf, Inf)$value  
  sigma <- min(sd(x), IQR(x) / 1.34)  
  (8 * sqrt(pi) * K / (3 * sigma2^2))^(1/5) * sigma * n^(-1/5)  
}
```




# Implementation of Kernel Density Estimate with Bandwidth Selection

```
dens_R1 <- function(x,
                    kernel,
                    bandwidth,
                    ...,
                    points = 512L) {
  if(is.function(bandwidth)) {
    bw <- bandwidth(x, kernel, ...)
  } else {
    bw <- bandwidth
  }
  p <- seq(min(x) - bw, max(x) + bw, length.out = points)
  y <- R_dens_for(x, p, kernel, bw)
  structure(
    list(
      x = p,
      y = y,
      bw = bw
    ),
    class = "dens"
  )
}
```

# Plotting the Density Estimate with Bandwidth Selection


The LOOCV bandwidth is 0.1155167 and the oracle bandwidth is 0.4746436.



# Profiling the Implementation

## Profiling the Kernel Density Estimation


```
dens_R1 <- function(x,
                    kernel,
                    bandwidth,
                    points = 512L) {
  if(is.function(bandwidth)) {
    bw <- bandwidth(x, kernel)
  } else {
    bw <- bandwidth
  }
  p <- seq(min(x) - bw, max(x) + bw, length.out = points)
  y <- R_dens_for(x, p, kernel, bw)
  structure(
    list(
      x = p,
      y = y,
      bw = bw
    ),
    class = "dens"
  )
}
```



# Profiling the Implementation

## Profiling the Bandwidth Selection

```
bw_cv_R2 <- function(x, kernel) {  
  cv_func <- function(l) {  
    if(l < .Machine$double.eps) Inf  
    n <- length(x)  
    K <- numeric(n)  
    for(i in 1:n) {  
      K[i] <- sum(kernel((x[i] - x[-i]) / l))  
    }  
    cv <- sum(log(K[K > .Machine$double.eps]))  
    n * log((n - 1) * l) - cv  
  }  
  suppressWarnings(optimize(cv_func, c(0, 2)))$minimum  
}
```




# Alternative Implementations of Calculating the Density Estimate

```
R_dens <- function(x, p, kernel, bandwidth) {  
  m <- length(p)  
  n <- length(x)  
  result <- numeric(m)  
  for(i in 1:m) {  
    result[i] <- sum(kernel((p[i] - x) / bandwidth))  
  }  
  result / (n * bandwidth)  
}  
  
R_dens_for <- function(x, p, kernel, bandwidth) {  
  m <- length(p)  
  n <- length(x)  
  result <- numeric(m)  
  for(i in 1:m) {  
    for(j in 1:n) {  
      result[i] <- result[i] + kernel((p[i] - x[j]) / bandwidth)  
    }  
  }  
  result / (n * bandwidth)  
}  
  
R_dens_apply <- function(x, p, kernel, bandwidth) {  
  app_func <- function(t) sum(kernel((t - x) / bandwidth))  
  result <- sapply(p, app_func)  
  result / (length(x) * bandwidth)  
}
```

# Benchmarking Alternative Implementations using bench Package

Plot



# Benchmarking Alternative Implementations using bench Package

Table

|   | expression              | median                | mem_alloc              |
|---|-------------------------|-----------------------|------------------------|
|   | <i>&lt;bch:expr&gt;</i> | <i>&lt;bch:tm&gt;</i> | <i>&lt;bch:byt&gt;</i> |
| 1 | ONE_FOR_LOOP            | 19.61ms               | 35.5MB                 |
| 2 | APPLY                   | 20.83ms               | 35.4MB                 |
| 3 | TWO_FOR_LOOP            | 1.03s                 | 79.1KB                 |

# Alternative LOOCV Implementations


Let  $K_{i,j} = K((x_i - x_j)/\lambda)$  and notice  $K_{i,j} = K_{j,i}$ .

```
bw_cv_R <- function(x, kernel, max_bw = 2) {  
  cv_func <- function(l) {  
    if(l < .Machine$double.eps) Inf  
    n <- length(x)  
    K <- numeric(n)  
    for(i in 2:n) {  
      index <- 1:(i - 1)  
      tmp <- kernel((x[i] - x[index]) / l)  
      K[i] <- K[i] + sum(tmp)  
      K[index] <- K[index] + tmp[index]  
    }  
    cv <- sum(log(K[K > .Machine$double.eps]))  
    n * log((n - 1) * l) - cv  
  }  
  suppressWarnings(optimize(cv_func, c(0, max_bw)))$minimum  
}
```



# Alternative LOOCV Implementation Benchmarks

Plot



# Alternative LOOCV Implementation Benchmarks

Table

|   | expression              | median                | mem_alloc              |
|---|-------------------------|-----------------------|------------------------|
|   | <i>&lt;bch:expr&gt;</i> | <i>&lt;bch:tm&gt;</i> | <i>&lt;bch:byt&gt;</i> |
| 1 | TRIANGLE                | 357ms                 | 433MB                  |
| 2 | RECTANGLE               | 407ms                 | 738MB                  |

# Improving Further with RCPP Package

## Implementing the kernel in RCPP

Implementing the kernel in RCPP could improve performance per the profiling.

```
// [[Rcpp::export]]  
NumericVector e_kernel_cpp(NumericVector x) {  
  int n = x.size();  
  NumericVector result(n);  
  for(int i = 0; i < n; ++i)  
    result[i] = std::abs(x[i]) <= 1 ? 0.75 * (1 - x[i] * x[i]) : 0;  
  return result;  
}
```

# Improving Further with RCPP Package

## Calling Kernel from R in RCPP

```
// [[Rcpp::export]]
NumericVector dens_rcpp_partial(NumericVector x,
                                Function kernel,
                                double bandwidth,
                                NumericVector points) {
  int n = x.size(), m = points.size();
  NumericVector result(m);
  for(int i = 0; i < m; ++i) {
    NumericVector call = kernel((points[i] - x) / bandwidth);
    result[i] = sum(call) / (n * bandwidth);
  }
  return result;
}
```

# Improving Further with RCPP Package

## Calling Kernel from R in RCPP

```
// [[Rcpp::export]]
double bw_cv_rcpp_partial(NumericVector x,
                          Function kernel,
                          double bandwidth) {

  int n = x.size();
  NumericVector K(n);
  double result;
  for(int i = 1; i < n; ++i) {
    Range r = Range(0, i - 1);
    NumericVector tmp = kernel((x[r] - x[i]) / bandwidth);
    for(int j = 0; j < i; ++j)
      K[j] += tmp[j], K[i] += tmp[j];
  }
  for(int s = 0; s < n; ++s)
    if(K[s] > std::numeric_limits<double>::min()) result += std::log(K[s]);
  return n * log((n - 1) * bandwidth) - result;
}
```

# Improving Further with RCPP Package

## Full RCPP Implementation

```
// [[Rcpp::export]]
NumericVector dens_rcpp(NumericVector x,
                        double bandwidth,
                        NumericVector points) {
  int n = x.size(), m = points.size();
  NumericVector result(m);
  for(int i = 0; i < m; ++i) {
    NumericVector call = e_kernel_cpp((points[i] - x) / bandwidth);
    result[i] = sum(call) / (n * bandwidth);
  }
  return result;
}
```


# Improving Further with RCPP Package

## Full RCPP Implementation

```
// [[Rcpp::export]]
double bw_cv_rcpp(NumericVector x,
                  double bandwidth) {
    int n = x.size();
    NumericVector K(n);
    double result;
    for(int i = 1; i < n; ++i) {
        Range r = Range(0, i - 1);
        NumericVector tmp = e_kernel_cpp((x[i] - x[r]) / bandwidth);
        for(int j = 0; j < i; ++j)
            K[j] += tmp[j], K[i] += tmp[j];
    }
    for(int s = 0; s < n; ++s)
        if(K[s] > std::numeric_limits<double>::min()) result += std::log(K[s]);
    return n * std::log((n - 1) * bandwidth) - result;
}
```

# Benchmarking RCPP Density Calculation

Plot






# Benchmarking RCPP Density Calculation

Table

|   | expression              | median                | mem_alloc              |
|---|-------------------------|-----------------------|------------------------|
|   | <i>&lt;bch:expr&gt;</i> | <i>&lt;bch:tm&gt;</i> | <i>&lt;bch:byt&gt;</i> |
| 1 | CPP_FULL                | 5.73ms                | 15.7MB                 |
| 2 | R_CPP_KERNEL            | 9.28ms                | 18.2MB                 |
| 3 | R                       | 13.62ms               | 35.5MB                 |
| 4 | CPP_CPP_KERNEL          | 23.66ms               | 18.2MB                 |
| 5 | CPP                     | 28.32ms               | 35.4MB                 |

# Benchmarking RCPP Bandwidth Selection

Plot



# Benchmarking RCPP Bandwidth Selection

Table

|   | expression              | median                | mem_alloc              |
|---|-------------------------|-----------------------|------------------------|
|   | <i>&lt;bch:expr&gt;</i> | <i>&lt;bch:tm&gt;</i> | <i>&lt;bch:byt&gt;</i> |
| 1 | CPP_FULL                | 5.73ms                | 15.7MB                 |
| 2 | R_CPP_KERNEL            | 9.28ms                | 18.2MB                 |
| 3 | R                       | 13.62ms               | 35.5MB                 |
| 4 | CPP_CPP_KERNEL          | 23.66ms               | 18.2MB                 |
| 5 | CPP                     | 28.32ms               | 35.4MB                 |

# Improving Further Using R's C API

## Implementation of Density Calculation

```
#include <Rinternals.h>
#include <R.h>

SEXP C_dens(SEXP x, SEXP p, SEXP kernel, SEXP bw, SEXP rho) {
    int n = length(x), m = length(p);
    SEXP dens = PROTECT(allocVector(REALSXP, m));
    SEXP tmp = PROTECT(allocVector(REALSXP, n));
    SEXP K_Call = PROTECT(lang2(kernel, R_NilValue));
    double *x_ = REAL(x), *p_ = REAL(p), *tmp_ = REAL(tmp), *dens_ = REAL(dens);
    double bw_ = REAL(bw)[0];

    memset(dens_, 0, sizeof(double) * m);

    for(int i = 0; i < m; ++i) {
        for(int j = 0; j < n; ++j)
            tmp_[j] = (p_[i] - x_[j]) / bw_;
        SETCADR(K_Call, tmp);
        SEXP result = eval(K_Call, rho);
        double *result_ = REAL(result);
        for(int j = 0; j < n; ++j)
            dens_[i] += result_[j];
        dens_[i] /= n * bw_;
    }
    UNPROTECT(3);
    return dens;
}
```

# Improving Further Using R's C API

## Implementation of LOOCV

```
#include <Rinternals.h>
#include <R.h>
#include <float.h>

SEXP C_cv(SEXP x, SEXP fn, SEXP lambda, SEXP rho) {
  if (REAL(lambda)[0] < DBL_EPSILON) return ScalarReal(INFINITY);
  int n = length(x);
  SEXP K = PROTECT(allocVector(REALSXP, n));
  SEXP fn_call = PROTECT(lang2(fn, R_NilValue));
  SEXP out = PROTECT(ScalarReal(0.0));
  double *x_ = REAL(x), *K_ = REAL(K), *out_ = REAL(out);
  double h_ = REAL(lambda)[0];
  memset(K_, 0, sizeof(double) * n);
  for (int i = 1; i < n; ++i) {
    SEXP tmp = PROTECT(allocVector(REALSXP, i));
    double *tmp_ = REAL(tmp);
    for (int k = 0; k < i; ++k)
      tmp_[k] = (x_[i] - x_[k]) / h_;
    SETCADR(fn_call, tmp);
    SEXP s = eval(fn_call, rho);
    double *s_ = REAL(s);
    UNPROTECT(1);
    for (int j = 0; j < i; ++j) {
      K_[i] += s_[j];
      K_[j] += s_[j];
    }
  }
  for (int i = 0; i < n; ++i)
    if (K_[i] > DBL_EPSILON) *out_ += log(K_[i]);
  *out_ = n * log((n - 1) * h_) - (*out_);
  UNPROTECT(3);
  return out;
}
```

# Improving Further Using R's C API


## Calling the C Code

Compile to shared library using R CMD SHLIB command. If used in a package a little more work is required. For now just used `dyn.load` function in R to link the shared library.

```
bw_cv <- function(x, kernel, max_bw = 2) {  
  cv_func <- function(l) .Call("C_cv", x, kernel, l, environment())  
  suppressWarnings(optimize(cv_func, c(0, max_bw)))$minimum  
}
```

# Benchmarking All Density Calculation Implementations

Plot



# Benchmarking All Density Calculation Implementations


Table

|   | expression              | median                | mem_alloc              |
|---|-------------------------|-----------------------|------------------------|
|   | <i>&lt;bch:expr&gt;</i> | <i>&lt;bch:tm&gt;</i> | <i>&lt;bch:byt&gt;</i> |
| 1 | CPP_FULL                | 5.81ms                | 15.7MB                 |
| 2 | C_CPP_KERNEL            | 7.42ms                | 10.4MB                 |
| 3 | R_CPP_KERNEL            | 9.39ms                | 18.2MB                 |
| 4 | C                       | 11.89ms               | 27.6MB                 |
| 5 | R                       | 13.74ms               | 35.5MB                 |
| 6 | CPP_CPP_KERNEL          | 23.9ms                | 18.2MB                 |
| 7 | CPP                     | 28.7ms                | 35.4MB                 |



# Benchmarking All LOOCV Implementations

Plot



# Benchmarking All LOOCV Implementations

Table

|   | expression              | median                | mem_alloc              |
|---|-------------------------|-----------------------|------------------------|
|   | <i>&lt;bch:expr&gt;</i> | <i>&lt;bch:tm&gt;</i> | <i>&lt;bch:byt&gt;</i> |
| 1 | CPP_FULL                | 103ms                 | 124MB                  |
| 2 | C_CPP_KERNEL            | 133ms                 | 163MB                  |
| 3 | C                       | 183ms                 | 279MB                  |
| 4 | R_CPP_KERNEL            | 320ms                 | 317MB                  |
| 5 | R                       | 363ms                 | 433MB                  |
| 6 | CPP_CPP_KERNEL          | 401ms                 | 162MB                  |
| 7 | CPP                     | 450ms                 | 278MB                  |

# Final implementation

```
dens <- function(x,
                 kernel,
                 bandwidth = bw_cv,
                 ...,
                 points = 512L) {
  if(is.function(bandwidth)) {
    bw <- bandwidth(x, kernel, ...)
  } else {
    bw <- bandwidth
  }
  p <- seq(min(x) - bw, max(x) + bw, length.out = points)
  y <- .Call("C_dens", x, p, kernel, bw, environment())
  structure(
    list(
      x = p,
      y = y,
      bw = bw
    ),
    class = "dens"
  )
}
```