

Adaptive Envelope Rejection Sampling

Applied to Log-concave Densities

Lucas Støjko Andersen

University of Copenhagen

November 8, 2022

Rejection Sampling

Let f, g be densities on \mathbb{R} with $\alpha f \leq g$ for $\alpha > 0$. Let U_1, U_2, \dots be i.i.d. with uniform distribution and Y_1, Y_2, \dots be i.i.d. with density g independent of the U_i 's. Define the stopping time

$$\sigma = \inf\{n \geq 1 \mid U_n \leq \alpha f(Y_n)/g(Y_n)\}. \quad (1)$$

Then Y_σ has density f . The densities need not be normalized, however, if they are then $\alpha \in (0, 1]$ and $1 - \alpha$ is the probability of rejection.

Rejection Sampling

Implementation

```
rejection_sampling <- function(n,
                                density,
                                env_density,
                                env_sampler,
                                alpha,
                                seed = NULL) {
  if(!is.null(seed)) set.seed(seed)
  samples <- numeric(n)
  succes <- tries <- 0
  for(s in 1:n) {
    reject <- TRUE
    while(reject) {
      tries <- tries + 1
      u0 <- runif(1)
      y0 <- env_sampler()
      env_y0 <- env_density(y0)
      dens_y0 <- density(y0)
      if(u0 <= alpha * dens_y0 / env_y0) {
        reject <- FALSE
        samples[s] <- y0
        succes <- succes + 1
      }
    }
  }
  list(samples, (tries - succes) / tries)
}
```

Rejection Sampling

Case Study

Let f be the density with

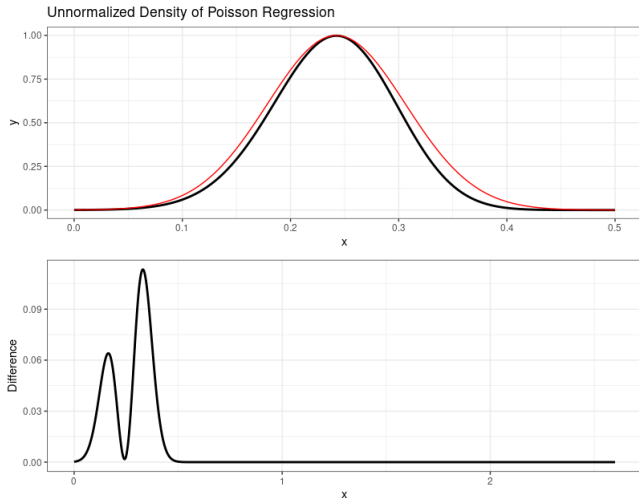
$$f(y) \propto p(y) = \prod_{i=1}^{100} \exp(yz_i x_i - e^{yx_i}). \quad (2)$$

```
poisreg <- function(x, z) {  
  force(x); force(z)  
  function(y) {  
    expyx <- sapply(y, function(s) sum(exp(s * x)))  
    exp(y * sum(x * z) - expyx)  
  }  
}  
  
poisreg_derv <- function(x, z) {  
  force(x)  
  force(z)  
  function(y) {  
    expyx <- sapply(y, function(s) sum(exp(s * x)))  
    x_expyx <- sapply(y, function(s) sum(x * exp(s * x)))  
    xz <- sum(x * z)  
    exp(y * xz - expyx) * (xz - x_expyx)  
  }  
}
```

Rejection Sampling

A Gaussian Envelope

The function $e^{(x-0.2423914)^2/(2 \cdot 0.004079805)}$ is an envelope of p with $\alpha = 1.351351 \cdot 10^{40}$.



Adaptive Rejection Sampling

```
adap_samp <- function(n, density, density_deriv, p, zb = c(-Inf, Inf), seed = NULL) {  
  if(!is.null(seed)) set.seed(seed)  
  p <- sort(unique(p))  
  densp <- density(p)  
  a <- density_deriv(p) / densp  
  b <- log(densp) - a * p  
  a_diff <- a[-length(a)] - a[-1]  
  check1 <- a[1] < 0 & zb[1] == -Inf  
  check2 <- a[length(a)] > 0 & zb[2] == Inf  
  if(check1 | check2)  
    stop("Envelope is not integrable. Choose different points.")  
  if(any(a == 0) | any(a_diff == 0))  
    stop("Division by zero. Choose different points.")  
  z <- c(zb[1], (b[-1] - b[-length(b)]) / a_diff, zb[2])  
  env_quantile <- get_env_quantile(a, b, z)  
  env_density <- get_env_density(a, b, z)  
  samples <- numeric(n)  
  succes <- tries <- 0  
  for(s in 1:n) {  
    reject <- TRUE  
    while(reject) {  
      tries <- tries + 1  
      u0 <- runif(2)  
      y0 <- env_quantile(u0[1])  
      env_y0 <- env_density(y0)  
      dens_y0 <- density(y0)  
      if(u0[2] <= dens_y0 / env_y0) {  
        reject <- FALSE  
        succes <- succes + 1  
        samples[s] <- y0  
      }  
    }  
  }  
  list(samples, (tries - succes) / tries)  
}
```

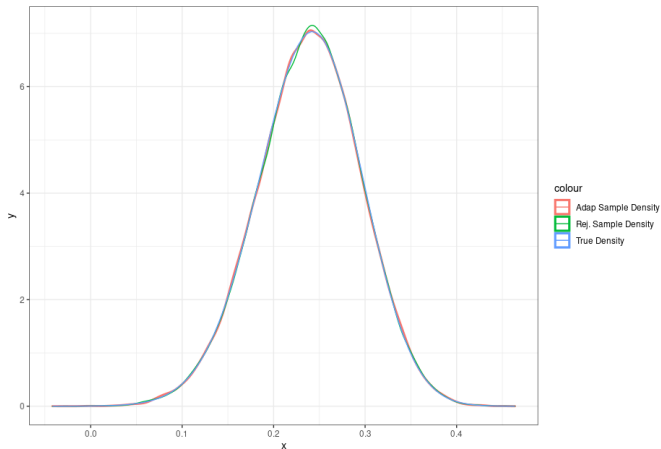
Adaptive Rejection Sampling - Continued

Implementation

```
get_env_quantile<- function(a, b, z) {  
  force(a); force(b); force(z)  
  az <- a * z[-length(z)]  
  R <- exp(b) * (exp(a * z[-1]) - exp(az)) / a  
  Q1 <- numeric(length(a) + 1)  
  Q1[2:length(Q1)] <- cumsum(R)  
  c <- Q1[length(Q1)]  
  function(q) {  
    ind <- c * q <= Q1  
    maxi <- which.max(ind) - 1  
    y <- c * q - Q1[maxi]  
    log(a[maxi] * y * exp(-b[maxi]) + exp(az[maxi])) / a[maxi]  
  }  
}  
  
get_env_density <- function(a, b, z) {  
  force(a); force(b); force(z)  
  function(x) {  
    if(x > z[length(z)] | x < z[1]) return(0)  
    maxi <- which.max(x <= z) - 1  
    exp(a[maxi] * x + b[maxi])  
  }  
}
```

Comparison of the Implementations

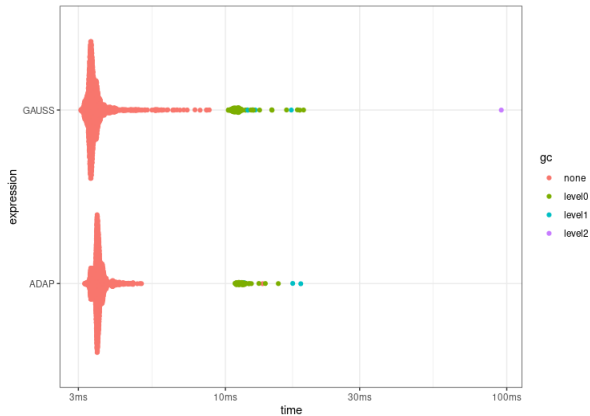
Simulating 100.000 samples with each implementation using the same seed. Using 0.15, 0.2, 0.28, 0.32 as the envelope points. The adaptive envelope had 0.09 rate of rejection and the gaussian envelope had 0.11 rate of rejection.



Comparison of the Implementations

Benchmarks Using `bench` Package

Benchmarking sampling 100 samples.



Comparison of the Implementations

Benchmarks Using `bench` Package

Benchmarking sampling 100 samples.

```
-----  
expression      min      median `itr/sec` mem_alloc `  
<bch:expr> <bch:tm> <bch:tm>      <dbl> <bch:byt>  
GAUSS           3.06ms    3.36ms    289.      928KB  
ADAP            3.16ms    3.52ms    281.      937KB  
|
```

Profiling the Implementation Using profvis Package

```

adap_samp <- function(n, density, density_deriv, p, zb = c(-
Inf, Inf), seed = NULL) {
  if(!is.null(seed)) set.seed(seed)
  p <- sort(unique(p))
  densp <- density(p)
  a <- density_deriv(p) / densp
  b <- log(densp) - a * p
  a_diff <- a[-length(a)] - a[-1]
  check1 <- a[1] > 0 & a[length(a)] < 0
  check2 <- a[length(a)] < 0 & is.finite( zb[1])
  check3 <- a[1] > 0 & is.finite(zb[2])
  if(!(check1 | check2 | check3))
    warning("Envelope is not integrable. Choose different
points.")
  if(any(a == 0) | any(a_diff == 0))
    stop("Division by zero. Choose different points.")
  z <- c(zb[1], (b[-1] - b[-length(b)]) / a_diff, zb[2])
  env_quantile <- get_env_quantile(a, b, z)
  env_density <- get_env_density(a, b, z)
  samples <- numeric(n)
  succes <- tries <- 0
  for(s in 1:n) {
    reject <- TRUE
    while(reject) {
      tries <- tries + 1
      u0 <- runif(2)
      y0 <- env_quantile(u0[1])
      env_y0 <- env_density(y0)
      dens_y0 <- density(y0)
      if(u0[2] <= dens_y0 / env_y0) {
        reject <- FALSE
        succes <- succes + 1
        samples[s] <- y0
      }
    }
  }
}

```

	0.4	20
	1.1	10
-51.5	46.5	350
-53.3	39.1	320
-72.1	43.6	400
-203.7	256.3	2170
	3.0	30
	1.7	10
	1.0	10

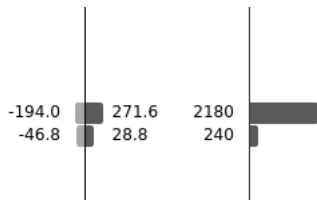
Profiling the Implementation Using profvis Package

```
get_env_quantile<- function(a, b, z) {  
  force(a); force(b); force(z)  
  az <- a * z[-length(z)]  
  R <- exp(b) * (exp(a * z[-1]) - exp(az)) / a  
  Q1 <- numeric(length(a) + 1)  
  Q1[2:length(Q1)] <- cumsum(R)  
  c <- Q1[length(Q1)]  
  function(q) {  
    ind <- c * q <= Q1  
    maxi <- which.max(ind) - 1  
    y <- c * q - Q1[maxi]  
    log(a[maxi] * y * exp(-b[maxi]) + exp(az[maxi])) / a[maxi]  
  }  
}  
  
get_env_density <- function(a, b, z) {  
  force(a); force(b); force(z)  
  function(x) {  
    if(x > z[length(z)] | x < z[1]) return(0)  
    maxi <- which.max(x <= z) - 1  
    exp(a[maxi] * x + b[maxi])  
  }  
}
```

	1.7	10
	7.1	50
	5.0	40
	3.8	30
-20.1	7.7	60
	2.2	20
-40.8	18.0	170
-10.6	15.2	120
-20.8	14.5	120

Profiling the Implementation Using profvis Package

```
poisreg <- function(x, z) {  
  force(x)  
  force(z)  
  function(y) {  
    expyx <- sapply(y, function(s) sum(exp(s * x)))  
    exp(y * sum(x * z) - expyx)  
  }  
}
```



Improving the Implementation

Optimizing use of runif

```
u_samples <- runif(2 * n); k_stop <- n; k <- 1
succes <- tries <- 0
for(s in 1:n) {
  reject <- TRUE
  while(reject) {
    tries <- tries + 1
    if(k == k_stop) {
      u_samples <- runif(2 * (n - (s - 1)))
      k_stop <- n - (s - 1) + 1
      k <- 1
    }
    u0 <- u_samples[2 * (k - 1) + 1]
    u1 <- u_samples[2 * (k - 1) + 2]
    k <- k + 1
    y0 <- env_quantile(u0)
    env_y0 <- env_density(y0)
    dens_y0 <- density(y0)
    if(u1 <= dens_y0 / env_y0) {
      reject <- FALSE
      samples[s] <- y0
      succes <- succes + 1
    }
  }
}
```

Improving the Implementation

Implementation of Envelope Quantile and Density Function in RCPP

```
// [[Rcpp::export]]
double RCPP_env_density(double x,
                        NumericVector a,
                        NumericVector b,
                        NumericVector z) {

  int m = z.size();
  if(x < z[0] || x > z[m - 1]) return 0;
  int maxi;
  for(maxi = 1; maxi < m; ++maxi)
    if(x <= z[maxi]) break;
  maxi -= 1;
  return std::exp(a[maxi] * x + b[maxi]);
}

// [[Rcpp::export]]
double RCPP_env_quantile(double x,
                        NumericVector a,
                        NumericVector b,
                        NumericVector z,
                        NumericVector az,
                        NumericVector Q) {

  int m = Q.size(), maxi;
  double c = Q[m - 1];
  for(maxi = 0; maxi < m; ++maxi)
    if(c * x <= Q[maxi]) break;
  maxi -= 1;
  double y = c * x - Q[maxi];
  return std::log(a[maxi] * y *
                 std::exp(-b[maxi]) + std::exp(az[maxi])) / a[maxi];
}
```

Improving the Implementation

Implementation of Envelope Quantile and Density Function in RCPP

```
get_env_quantile_cpp<- function(a, b, z) {  
  force(a); force(b); force(z)  
  az <- a * z[-length(z)]  
  R <- exp(b) * (exp(a * z[-1]) - exp(az)) / a  
  Q1 <- numeric(length(a) + 1)  
  Q1[2:length(Q1)] <- cumsum(R)  
  c <- Q1[length(Q1)]  
  function(q) {  
    RCPP_env_quantile(q, a, b, z, az, Q1)  
  }  
}  
  
get_env_density_cpp <- function(a, b, z) {  
  force(a); force(b); force(z)  
  function(x) {  
    RCPP_env_density(x, a, b, z)  
  }  
}
```


Improving the Implementation

Benchmarking

expression	min	median	itr/s... ¹	mem_a... ² g
<bch:expr>	<bch:tm>	<bch:>	<dbl>	<bch:b>
GAUSS	3.11ms	3.34ms	281.	794KB
ADAP	3.28ms	3.52ms	275.	488KB
ADAP_STR	2.94ms	3.16ms	313.	216KB
ADAP_STR_CPP	3.64ms	3.83ms	257.	794KB

Improving the Implementation

Partial RCPP Implementation

```
// [[Rcpp::export]]
List RCPP_adap_samp_partial(int n,
                           Function density,
                           NumericVector a,
                           NumericVector b,
                           NumericVector z) {

  int m = a.size();
  std::vector<double> Q(m + 1);
  std::vector<double> az(m);
  Q[0] = 0;
  for(int i = 1; i < m + 1; ++i) {
    az[i - 1] = a[i - 1] * z[i - 1];
    Q[i] = Q[i - 1] +
      std::exp(b[i - 1]) *
      (std::exp(a[i - 1] * z[i]) - std::exp(az[i - 1])) / a[i - 1];
  }
}
```

```
NumericVector samples(n);
int accepts = 0, tries = 0;
for(int i = 0; i < n; ++i) {
  int reject = 1;
  while(reject == 1) {
    ++tries;
    double u0 = R::runif(0, 1);
    double u1 = R::runif(0, 1);
    double y0 = env_quantile(u0, a, b, az, Q);
    double env_y0 = env_density(y0, a, b, z);
    NumericVector dens_y0 = density(y0);
    if(u1 <= dens_y0[0] / env_y0) {
      reject = 0;
      samples[i] = y0;
      ++accepts;
    }
  }
}
NumericVector rate(1);
rate[0] = ((double) tries - (double) accepts) / (double) tries;
return List::create(samples, rate);
}
```

Improving the Implementation

Partial RCPP Implementation

```
double env_quantile(double x,
                    NumericVector &a,
                    NumericVector &b,
                    std::vector<double> &az,
                    std::vector<double> &Q) {
    int n = a.size();
    double c = Q[n];
    int maxi;
    for(maxi = 0; maxi < n + 1; ++maxi)
        if(c * x <= Q[maxi]) break;
    maxi -= 1;
    double y = c * x - Q[maxi];
    return std::log(a[maxi] * y * std::exp(-b[maxi]) +
                   std::exp(az[maxi])) / a[maxi];
}

double env_density(double x,
                   NumericVector &a,
                   NumericVector &b,
                   NumericVector &z) {
    int n = a.size();
    if(x > z[n] || x < z[0])
        return 0;
    int maxi;
    for(maxi = 0; maxi < n + 1; ++maxi)
        if(x <= z[maxi]) break;
    maxi -= 1;
    return std::exp(a[maxi] * x + b[maxi]);
}
```

Improving the Implementation

Partial RCPP Implementation

```
adap_samp_cpp_partial <- function(n,
                                   density,
                                   density_deriv,
                                   p,
                                   zb = c(-Inf, Inf),
                                   seed = NULL) {
  if(!is.null(seed)) set.seed(seed)
  p <- sort(unique(p))
  densp <- density(p)
  a <- density_deriv(p) / densp
  b <- log(densp) - a * p
  a_diff <- a[-length(a)] - a[-1]
  check1 <- a[1] < 0 & zb[1] == -Inf
  check2 <- a[length(a)] > 0 & zb[2] == Inf
  if(check1 | check2)
    stop("Envelope is not integrable. Choose different points.")
  if(any(a == 0) | any(a_diff == 0))
    stop("Divison by zero. Choose different points.")
  z <- c(zb[1], (b[-1] - b[-length(b)]) / a_diff, zb[2])

  RCPP_adap_samp_partial(n, density, a, b ,z)
}
```

Improving the Implementation

Benchmark

expression	min	median	`itr/sec`	mem_al... ¹
<bch:expr>	<bch:tm>	<bch:tm>	<dbl>	<bch:by>
1 GAUSS	3.12ms	3.31ms	293.	794KB
2 ADAP	3.28ms	3.53ms	280.	488KB
3 ADAP_STR	2.95ms	3.16ms	314.	216KB
4 ADAP_STR_CPP	3.63ms	3.82ms	257.	784KB
5 ADAP_PARTIAL_CPP	4.46ms	4.74ms	208.	401KB

Implementing Adaptive Envelope Sampling Using R's C API

```
#include <R.h>
#include <math.h>
#include <Rinternals.h>
#include <R_ext/Random.h> // ACCESS TO UNIF NUMBER GENERATOR
```

```
SEXP C_adap_samp(SEXP n,
                 SEXP density,
                 SEXP a,
                 SEXP b,
                 SEXP z,
                 SEXP rho) {
    int m = length(a);
    int N = INTEGER(n)[0];
    double *Q = (double *)malloc(sizeof(double) * (m + 1));
    double *az = (double *)malloc(sizeof(double) * m);
    double *a_ = REAL(a), *b_ = REAL(b), *z_ = REAL(z);

    Q[0] = 0;
    for(int i = 1; i < m + 1; ++i) {
        az[i - 1] = a_[i - 1] * z_[i - 1];
        Q[i] = Q[i - 1] +
            exp(b_[i - 1]) * (exp(a_[i - 1] * z_[i]) - exp(az[i - 1])) /
            a_[i - 1];
    }
}
```

```
SEXP density_call = PROTECT(lang2(density, R_NilValue));
SEXP samples = PROTECT(allocVector(REALSXP, N));
double *samples_ = REAL(samples);
int accepts = 0, tries = 0;
GetRNGstate();
for(int i = 0; i < N; ++i) {
    int reject = 1;
    while(reject == 1) {
        ++tries;
        double u0 = unif_rand();
        double u1 = unif_rand();
        double y0 = env_quantile(u0, a_, b_, az, Q, m);
        double env_y0 = env_density(y0, a_, b_, z_, m);
        SETCADR(density_call, PROTECT(ScalarReal(y0)));
        SEXP dens_y0 = eval(density_call, rho);
        UNPROTECT(1);
        if(u1 <= REAL(dens_y0)[0] / env_y0) {
            reject = 0;
            samples_[i] = y0;
            ++accepts;
        }
    }
}
PutRNGstate();
SEXP values = PROTECT(allocVector(VECSXP, 2));
double rate = ((double) tries - (double) accepts) / (double) tries;
SET_VECTOR_ELT(values, 0, samples);
SET_VECTOR_ELT(values, 1, ScalarReal(rate));
UNPROTECT(3);
free(Q);
free(az);
return values;
}
```

Implementing Adaptive Envelope Sampling Using R's C API

```
double env_quantile(double x,  
                   double *a,  
                   double *b,  
                   double *az,  
                   double *Q,  
                   int n) {  
  
    double c = Q[n];  
    int maxi;  
    for(maxi = 0; maxi < n + 1; ++maxi)  
        if(c * x <= Q[maxi]) break;  
    maxi -= 1;  
    double y = c * x - Q[maxi];  
    return log(a[maxi] * y * exp(-b[maxi]) + exp(az[maxi])) / a[maxi];  
}
```

```
double env_density(double x,  
                  double *a,  
                  double *b,  
                  double *z,  
                  int n) {  
  
    if(x > z[n] || x < z[0])  
        return 0;  
    int maxi;  
    for(maxi = 0; maxi < n + 1; ++maxi)  
        if(x <= z[maxi]) break;  
    maxi -= 1;  
    return exp(a[maxi] * x + b[maxi]);  
}
```

Implementing Adaptive Envelope Sampling Using R's C API

```
adap_samp_c <- function(n, density, density_deriv, p, seed = NULL, zb = c(-Inf, Inf)) {  
  if(!is.null(seed)) set.seed(seed)  
  p <- sort(unique(p))  
  densp <- density(p)  
  a <- density_deriv(p) / densp  
  b <- log(densp) - a * p  
  a_diff <- a[-length(a)] - a[-1]  
  check1 <- a[1] < 0 & zb[1] == -Inf  
  check2 <- a[length(a)] > 0 & zb[2] == Inf  
  if(check1 | check2)  
    stop("Envelope is not integrable. Choose different points.")  
  if(any(a == 0) | any(a_diff == 0))  
    stop("Division by zero. Choose different points.")  
  z <- c(zb[1], (b[-1] - b[-length(b)]) / a_diff, zb[2])  
  
  .Call("C_adap_samp",  
        as.integer(n),  
        density,  
        a,  
        b,  
        z,  
        environment())  
}
```

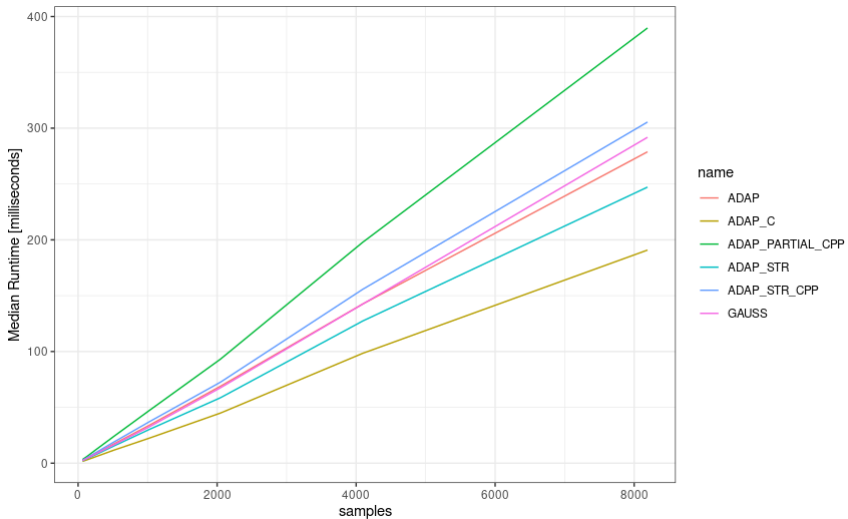

Final Benchmarks

Table

expression	min	median	`itr/sec`	mem_al... ¹
<i><bch:expr></i>	<i><bch:tm></i>	<i><bch:tm></i>	<i><dbl></i>	<i><bch:by></i>
GAUSS	3.12ms	3.32ms	293.	794KB
ADAP	3.28ms	3.52ms	281.	488KB
ADAP_STR	2.93ms	3.17ms	313.	216KB
ADAP_STR_CPP	3.64ms	3.82ms	257.	784KB
ADAP_PARTIAL_CPP	4.49ms	4.75ms	208.	206KB
ADAP_C	2.23ms	2.38ms	413.	407KB

Final Benchmarks

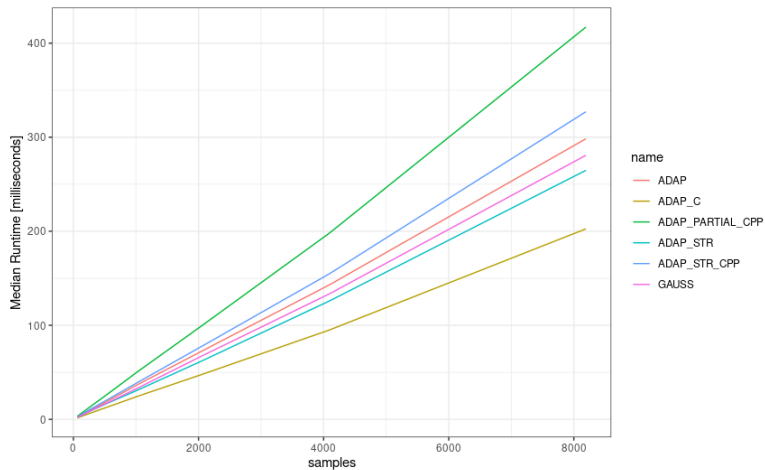
Scaling



Final Benchmarks

Scaling

Different points 0.1, 0.2, 0.3 for adaptive envelope.



Density of Samples from C Implementation

